

Gachix:

A binary cache for Nix over Git

Bachelor Thesis

Department of Computer Science
Department of Mathematics and Computer Science
Computer Networks Group
cn.dmi.unibas.ch

Examiner: Prof. Dr. Christian Tschudin
Supervisor: Dr. Erick Lavoie

Ephraim Siegfried
e.siegfried@unibas.ch

December 17, 2025

Abstract

This project develops a binary cache for Nix packages using Git's content-addressable filesystem. This approach improves upon traditional input-addressable packages by reducing memory usage and enhancing trust. Leveraging Git provides a key advantage: simple peer-to-peer replication of the binary cache across multiple nodes. The core work involves modeling package dependency graphs and user profiles within Git and creating an interface for Nix to interact with this system. The project will be evaluated through performance benchmarks measuring memory usage and package retrieval speed, alongside functional tests of the peer-to-peer replication and Nix interface.

Table of Contents

1	Introduction	5
2	Background	6
2.1.	Nix	6
2.1.1.	Nix Store	6
2.1.2.	Deployment Pipeline	7
2.1.3.	Binary Cache Interface	8
2.1.4.	Daemon Protocol	9
2.2.	Git	10
2.2.1.	Objects	10
2.2.2.	Replication	11
3	Design	12
3.1.	Mapping Nix to Git	12
3.1.1.	Packages	13
3.1.2.	Dependency Management	14
3.1.3.	User Profiles	15
3.2.	Cache Interface	15
3.2.1.	Narinfo endpoint	15
3.2.2.	NAR endpoint	16
3.3.	Replication Protocol	16
3.3.1.	Constructing Package Closures	17
3.3.2.	Fetching Packages from Replicas	18
3.3.3.	Nix Daemons	19
4	Implementation	20
4.1.	Architecture	20
4.2.	Concurrency	22
4.3.	Nix Archive	23
4.3.1.	Simple Encoder	23
4.3.2.	Stream Encoder	24

4.4. Nix Daemon	24
4.4.1. Libraries	25
4.5. Content and Input Addressing Schemes	25
4.6. Limitations	26
5 Evaluation	27
5.1. Functional Comparison to other Cache implementations	27
5.2. Package Retrieval Latency	28
5.2.1. Methodology	28
5.2.2. Result	28
5.2.3. Discussion	30
5.3. Package Storage	30
5.3.1. Methodology	30
5.3.2. Result	31
5.3.3. Discussion	31
5.4. Deployment on any Unix Machine	31
5.4.1. Methodology	31
5.4.2. Result	31
5.4.3. Discussion	32
5.5. Nix Transparency	32
5.5.1. Methodology	32
5.5.2. Result	32
5.5.3. Discussion	32
6 Related Work	33
6.1. Snix	33
6.2. Laut	33
7 Conclusion	34

1

Introduction

2

Background

2.1. Nix

Nix is a declarative and purely functional package manager founded by Eelco Dolstra. [1] The main purpose of Nix is to solve reproducibility shortcomings of other package managers. When packaging an application with RPM for example, the developer is supposed to declare all the dependencies, but there is no guarantee that the declaration is complete. There might be a shared library provided by the local environment which is a dependency the developer does not know of. The issue is that the app will build and work correctly on the machine of the developer but might fail on the end user's machine. To solve this issue, Nix provides a functional language with which the developer can declaratively define all software components a package needs. Nix then ensures that these dependency specifications are complete and that Nix expressions are deterministic, i.e. building a Nix expression twice yields the same result. [2]

2.1.1. Nix Store

The Nix store is a read-only directory (usually located at `/nix/store`) where Nix stores objects. Store objects are source files, build artifacts (e.g. binaries) and derivations among other things. These objects can refer to other objects (e.g. dependencies in binaries). To prevent ambiguity, every object has a unique identifier, which is a hash. This hash is reflected in the path of the object, which is constructed as `/nix/store/<nix-hash>-<object-name>-<object-version>`. [3]

There are two primary methods for computing the hash. The standard approach is to hash the derivation associated with the store object. A derivation is a file specifying the build

recipe and the complete graph of dependencies; essentially, this hashes the inputs. [4] There is also an experimental feature where the hash is computed directly from the final contents of the built object itself. [5] The former scheme is known as input-addressing, and the latter as content-addressing.

2.1.2. Deployment Pipeline

To produce a package in Nix, a derivation has to be produced. A derivation is a build plan that specifies how to create one or more output objects in the Nix store (it has the `.drv` extension). It also pins down the run and build time dependencies and specifies what the path of the output will be. It is an intermediate artifact generated when a Nix package expression is evaluated, analogous to an object file (`*.o`) in a C compilation process.

To build a derivation, Nix first ensures all dependent derivations are built. It then runs the builder in an isolated sandbox. In the sandbox only exclusively declared build and runtime dependencies can be accessed (e.g. `/bin` gets pruned from the environment variable `PATH`) and network access is limited. This makes component builders pure; when the deployer fails to specify a dependency explicitly, the component will fail deterministically. [1] The result of the build process is an object in the Nix store. Subsequently, we denote the output objects as artifacts.

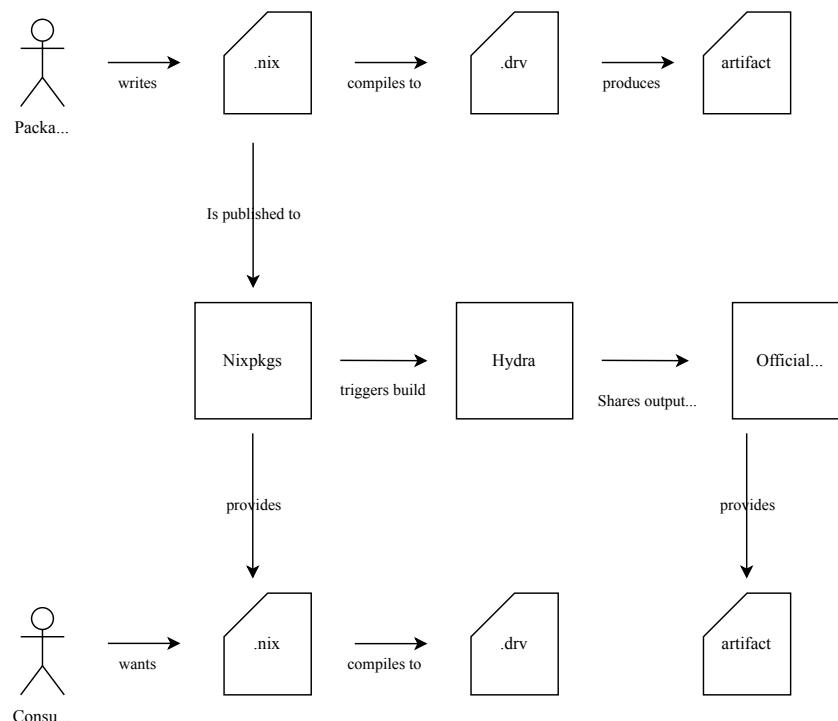


Figure 1: Nix Deployment Pipeline

If an author wants to share a package with others, the author needs to put the Nix expression which produces the artifact in a public registry. The official Nix registry called *Nixpkgs* is the place where most packages get published. It is maintained as a Git repository stored on GitHub. To add a package there, the author needs to make a pull request with the new package expression to be added. The expression gets reviewed by a trusted set of community members. Once accepted, the Nix expression will be added to the registry.

Users can build packages by specifying the registry and the name of the package. Nix will download the expression from the registry and produce a derivation. The derivation specifies the path of the artifact in the Nix store. To avoid building the artifacts locally, which can take a long time, users can benefit from binary caches, which are called substituers in Nix. With the default installation of Nix, there is only one substituer which is the official binary cache.¹ This cache has most artifacts which are in the official Nixpkgs registry. These official artifacts are built on Hydra, which is a continuous build system. Using the package identifier retrieved from the derivation, Nix will fetch the package from the binary cache. Package authors can also publish Nix expressions on a private registry and publish artifacts on a custom binary cache such as Cachix, which is a platform which both hosts and manages binary caches.². The benefit of publishing on the Nixpkgs registry is that artifacts will be made available at the official cache, which is set as trusted and available in every Nix installation.

2.1.3. Binary Cache Interface

The Nix binary cache interface provides a set of endpoints with which a user can retrieve metadata of packages and package contents.

Package contents are served in the NAR archive format [6]. This format closely follows the abstract specification of a file system object tree.

TODO: Improve explanation of NAR

Package metadata is served in the Narinfo format. It is a key value mapping with the following keys:

- **StorePath:** The full store path
- **URL:** The URL of the NAR fetching endpoint, relative to the binary cache
- **Compression:** The compression format of the served NAR
- **FileHash:** The hash of the compressed NAR

¹<https://cache.nixos.org/>

²<https://www.cachix.org>

- **FileSize**: The size of the compressed NAR
- **NarHash**: The hash of the NAR
- **NarSize**: The size of the NAR
- **Deriver**: The derivation which specifies the store object. It is the basename of the Nix path.
- **System**: The platform type of this binary, if known.
- **References**: A set of store paths which are direct runtime dependencies, separated by whitespace.
- **Sig**: A signature over the StorePath, NarHash, NarSize, and references fields usning ED25519 public-key signature system.

The most important endpoints of the binary cache API are:

- GET /<nix-hash>.narinfo: Retrieves the Narinfo for a given Nix hash, i.e. the hash substring in the path of a package.
- HEAD /<nix-hash>.narinfo: Used to check whether a package exists.
- GET /<url-in-narinfo>: Returns the compressed NAR. This endpoint is dependent on the endpoint given in the URL section of the Narinfo. Commonly, the endpoint is formed as GET /nar/<nix-hash>.nar.<compression>. [7]

2.1.4. Daemon Protocol

The Nix daemon is a service which runs runs Nix specific operations on behalf of non-root users. Most of the operations it can execute, can also be run via the Nix command line interface (CLI).

The main purpose of the Nix daemon is for multi-user Nix installations. In this mode the Nix store is owned by some privileged user to prevent other users to manipulate the Nix store in a malicious way (e.g. install a Trojan horse). When users use the Nix CLI, control is forwarded to the Nix daemon which is run under the owner of the Nix store. In this scenario, commands to the Nix daemon are dispatched through interprocess communication either via the socket (located at /nix/var/nix/daemon-socket/socket). [8]

There is also another usage of the Nix daemon. Nix is able to connect to remote Nix machines and perform operations on them. The initializer can for example order builds (`nix build --builders <remote-url>`) on the remote machine and fetch packages from the remote store (`nix copy --from <remote-url> <nix-path>`). Internally, Nix does this by connecting to the remote machine via SSH and starting the Nix daemon with `nix daemon --stdio`. This command starts a Nix daemon on the remote

machine and makes the remote daemon listen on standard I/O. Subsequently, both nodes communicate via the Nix daemon protocol. [9]

The Nix protocol starts with a handshake, where both parties agree on the protocol version they will use. They also exchange some configuration options. Subsequently there are 47 operations the parties can issue by sending the corresponding Opcode. [10]

2.2. Git

The official Git documentation advertises Git as a distributed version control system. More abstractly, it is a tool for manipulating a directed acyclic graph of objects and replicating these objects across repositories.

TODO: many more applications can be supported and those which can be supported have this dag property

2.2.1. Objects

Git objects are immutable and can only be added to the object database. There are four types of objects in Git:

- **Blob (Binary Large Object):** A blob is a sequence of bytes. It is used to store file data. The metadata of the blob, e.g. whether the file is executable or a symlink, is stored in the object that points to the blob, i.e. a tree.
- **Tree:** Is a collection of pointers to trees or blobs. It associates a name and other metadata with each pointer.
- **Commit:** Is a record which points to exactly one tree. It also contains the author of the commit, the time it was constructed and it can point to other commits which are called parents.
- **Reference:** Is a pointer to a Git object. **Direct References** can point to blobs, trees and commits; they supersede what are commonly known as references and tags. **Symbolic References** point to direct references.

Blobs, trees and commits are compressed and stored in the `.git/objects` directory. All objects in this directory are content addressed, i.e. they are identified by the hash of their contents. References are identified by their chosen name and are stored in `.git/refs`. [11]

2.2.2. Replication

Git can replicate objects across multiple repositories. One of the replication protocols is the *fetch* operation. With this operation, a repository can request and download objects from another repository, which are called remotes. The connection to remotes happens via HTTP or SSH.

To specify which objects should be downloaded, a *refspec* can be used. The *refspec* specifies which remote references should be downloaded to the local repository and how the received references should be named. This operation also downloads all objects which are reachable from the references specified. The *refspec* is a string which is structured as `<remote_references>:<local_references>`. For example, the refspec `refs/foo:/refs/bar` copies the remote reference `refs/foo` and names it `refs/bar` locally and downloads all objects reachable from `refs/foo`. It is also possible to specify multiple references by using globs, e.g. the reference `refs/heads/*` specifies all references which are in the namespace `refs/heads`.

3

Design

3.1. Mapping Nix to Git

This chapter introduces how Nix concepts are mapped to the Git object model. Figure 2 displays a simple Nix store with a package called *foo* which depends on packages *libfoo* and *bar*. This section shows which transformations are taken to have an equivalent model in Git which is represented in Figure 3.

Just like the Git object database, the Nix Store is an immutable collection of data. Apart from packages the Nix store also stores source files, derivations, links and other types of objects. Since a binary cache only serves binary packages, we can focus on only those type of objects.

Every top-level entry in the Nix store is uniquely addressable by some hash. There exist also files which are not uniquely addressable, which are inside those directories (e.g. a directory called ‘bin’). When mapping Nix entries to Git objects we have to make sure that the corresponding objects in Git are also uniquely addressable. Fortunately, this is already the case in Git, because Git uses the hash of the data content to address it.

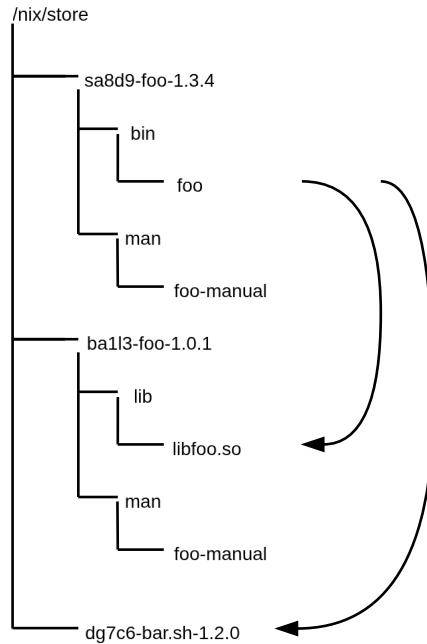


Figure 2: Nix Store

3.1.1. Packages

A Nix package is either a single file executable or a directory containing an executable. In Git, files can be mapped to blobs. This mapping will lose file metadata information: When constructing a blob from a file, and then reconstructing the file, it cannot be known whether the file is executable. There are many ways to solve this issue. One way to solve this, is to always assume that top-level files are executable, because the binary cache only serves executables. When sending the NAR of a top-level file to a client, we can always mark it as executable. Another way to solve this, is to construct a Git tree with the blob of the file as the single entry. In this tree, we can mark the blob as executable. In order to distinguish this tree from other trees, the blob in this special tree can be named with a unique magic value which marks the tree and the blob as a single file executable. The latter approach is favourable, because as presented in Section 3.1.2. we will create commits for each package, and commits can only point to trees and not blobs.

Package directories can be mapped to Git trees. These directories can contain symbolic links, files and other directories. Directories can be mapped to trees, symbolic links and files can be mapped to blobs, while metainfo such as the name of the objects can be stored in the trees.

Notice that after mapping Nix store entries to Git objects, the size of the database will have decreased. One reason is that Git compresses its objects. Another reason is that all objects

in Git are content addressed. If two package directories in the Nix store contain a file with the exact same content, Git will store this file as a blob only once. The package directories will be mapped to trees which point to the same blob.

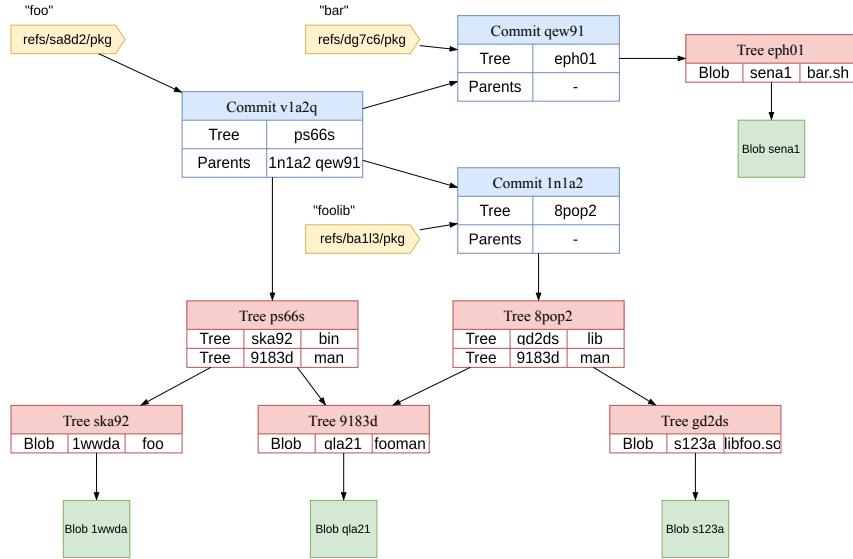


Figure 3: Gachix Git object model after transforming Nix store in Figure 2

3.1.2. Dependency Management

In order to track which packages have which runtime dependencies, Nix manages a Sqlite database. This database tracks which Nix paths have references to other paths in the Nix store. This information helps Nix to copy package closures to other stores and it prevents deleting packages which have references to them. In Gachix this dependency management is achieved using commit objects. For each package, a commit object is created where the commit tree is the tree containing the package contents as described above. The parents of the commit are the dependencies of the package which are also represented as commits. To find out what the dependency closure of a package called ‘foo’ is, we can recursively follow the parent pointers which is equivalent to running `git log <foo-commit-hash>`. To enable easy replication, we’ll need to ensure that every package is associated with exactly one commit hash (more in Section 3.3.). To ensure this, for every commit on each replica the commit message, the time stamp and the author field of the commit are set to constant values.

As a binary cache, Gachix needs to locate the corresponding commit hash given a Nix package hash, because the Nix binary cache interface requires that Nix packages are requested by the Nix hash. To maintain a mapping between Nix hashes and commit hashes I propose two solutions. One solution is to create a Git tree which serves as an index.

This special tree has all package commits as entries where the name of each entry is the respective Nix hash. With this special tree we can quickly lookup a package. A downside of this approach is that for each new entry the tree has to be copied (and the new entry be appended), because we cannot alter objects in the Git database. Another solution is to create a Git reference for each package which points to the corresponding commit object. The name of each reference contains the Nix hash. This approach also allows fast lookups. It is faster and requires less space since no objects have to be copied after package deletion or insertion. If we want to delete a package, we only have to remove the reference and call the garbage collector and Git will remove all objects associated with that package since these objects won't be reachable from a reference anymore. Because of these positive properties, Gachix uses the second approach.

3.1.3. User Profiles

Although this is not used in Gachix, an equivalent of Nix profiles and user environments can be achieved with the use of Git references. We can create a reference namespace (let's say for example `refs/myprofile`) containing symbolic references to direct references (e.g. `refs/myprofile/foo` pointing to `refs/<foo-nix-hash>/pkg`) which point to package commits. The symbolic references in this namespace point to the packages the user wishes to have in the environment. We can then create a worktree by merging all commits reachable from this reference namespace. This worktree is identical to the contents of `/nix/store` except that only active packages are contained in it.

3.2. Cache Interface

To integrate Gachix into the existing Nix ecosystem as a binary cache, it is easiest to implement the same API that the existing binary caches provide (see Section 2.1.3.). With this interface, a Nix user can add the URL of the Gachix service to the set of `substitutors`. Everytime the user adds a package, Nix will ask Gachix for package availability and fetch it if it is available.

3.2.1. Narinfo endpoint

The binary cache needs to serve metadata about the packages which is called Narinfo in Nix. Usually in other cache implementations the narinfo is computed on demand. In Gachix everytime a package is added to Gachix, the Narinfo is computed once and it is stored as a blob. Additionally, to associate this blob with a package, a reference is added under

`refs/<package-nix-hash>/narinfo` which points to this blob directly. Everytime a Narinfo is requested (with `GET /<package-nix-hash>.narinfo`) for a given Nix hash, Gachix serves the contents of this blob which it gets through this reference.

With the request `HEAD /<package-nix-hash>.nar` a client can ask the cache if it has the corresponding package. This request is handled by checking whether a reference exists containing the requested hash.

3.2.2. NAR endpoint

The Narinfo contains an URL, under which the NAR of a package can be downloaded. This link usually contains the hash of the NAR of the package. But because this endpoint can be chosen arbitrarily, I decided to put the hash of the package tree instead. It would also be possible to put the hash of the package commit, which would enable sending the whole package closure to the client. But since Nix only expects the contents of the requested package and not its dependencies, it is enough and faster to include the tree hash in the URL. The Nix user can then request a package with `GET refs/nar/\<tree-hash\>.nar`. Gachix then directly accesses the tree, converts it to a NAR and sends the archive to the client. The user can verify whether the hash corresponds to the package by unpacking the received NAR, converting it to a Git tree and comparing the tree hash with the hash in the URL.

3.3. Replication Protocol

This section explains how packages are added to the cache and replicated across peers. Gachix itself does not build packages and relies on external services. Gachix can communicate to the local Nix daemon, to remote Nix daemons and to remote Gachix peers (i.e. Git repositories managed by Gachix). With the Nix daemon protocol, Gachix can request metadata of store paths (e.g. runtime dependencies) and retrieve a package from the Nix store in the NAR format. With Gachix peers, Gachix uses the Git protocol to replicate commits and to fetch whole package dependency closures.

There is no policy yet on when Gachix adds packages to its repository. In the current version, it has to be manually run in the command line. A possible policy would be to always try to add a package when a Nix user requests one via the HTTP interface and it does not exist on the local repository. Gachix should then fetch the package from trusted replicas or may build it using the daemon protocol.

3.3.1. Constructing Package Closures

The current algorithm of adding a package closure is displayed in Algorithm 1. The algorithm receives the Nix store hash as an argument and returns the commit id associated with that package. It tries to recursively add package contents to the Git repository. At its core, it is similar to a depth first search algorithm (See lines 19-28 in Algorithm 1). The algorithm iterates through a package's direct dependencies (line 21). For each dependency d , it makes a recursive call to $\text{add_closure}(d)$ (line 23). This means it fully processes one dependency, including all its dependencies, before moving to the next direct dependency in the list. Once it has collected all commit hashes of the dependencies, it constructs a commit using the hash collection as parent commits (line 26).

Algorithm 1: Add package closure

```

1: procedure add_closure(path)
2:   if package_exists(nix_hash) then
3:     return commit_oid(nix_hash)
4:   end
5:    $\triangleright$  Ask gachix peers if they have already replicated the package closure
6:   for  $p \in P$  do
7:     if has_package(p, nix_hash) then
8:       return fetch_closure(nix_hash)
9:     end
10:    end
11:     $\triangleright$  Ask Nix daemons if they can provide package contents
12:    for  $d \in D$  do
13:      if has_package(d, nix_hash) then
14:        package  $\leftarrow$  fetch_package(nix_hash)
15:        break
16:      end
17:    end
18:     $\triangleright$  Get the nix hash of all dependent packages
19:    dependencies  $\leftarrow$  package.dependencies()
20:    parents  $\leftarrow \{\}$ 
21:    for  $d \in \text{dependencies}$  do
22:       $\triangleright$  Recursively fetch all commit oids
23:      dep_commit_oid  $\leftarrow$  add_closure(d)
24:      parents  $\cup \{\text{dep\_commit\_oid}\}$ 
25:    end
26:    commit_oid  $\leftarrow$  commit(package.tree, parents)
27:    add_reference(path, commit_oid)
28:    return commit_oid
29:  end

```

There are three base cases of the recursive algorithm. The algorithm does not recurse if it finds a leaf in the package dependency tree, i.e. a package without dependencies. It does

also not recurse if a package already exists in the local repository (lines 2-4). Another base case is when the package can be retrieved from peer replicas (lines 6-10).

3.3.2. Fetching Packages from Replicas

This section explains what the *fetch_package* function does in Algorithm 1 on line 8. Its main task is to fetch commits and references associated with a package. I will present two algorithms which perform this task. The first one is short and fast, but does not work because of an upstream Git bug. The second will be the one which is used in the current implementation.

We can fetch packages by using refsspecs (see Section 2.2.2.). With the refspec *refs/<nix-hash>/*:refs/<nix-hash>/** we specify that we want all references under the remote Nix hash namespace. Each such namespace has a *pkg* and a *narinfo* reference (see Section 3.1.2.). Git will fetch the blob from the *narinfo* reference and it will fetch all commits reachable from *pkg*. With this we will have all package contents from the whole package closure. What we will not have are the references to the dependencies of the package. For example, if a package *foo* has the dependency *bar* and we fetch with `git fetch peer refs/<nix-foo-hash>/*:refs/<foo-nix-hash>/*` we will receive the package contents of *foo* and *bar* but not the reference namespace *refs/<bar-nix-hash>*. To also get references from dependencies, we could add the namespace *refs/<nix-hash>/deps* containing symbolic references to all dependent packages. In the example above, when a peer constructs the package *foo*, it should also add the symbolic references *refs/<foo-nix-hash>/deps/<bar-nix-hash>/pkg* and *refs/<foo-nix-hash>/deps/<bar-nix-hash>/narinfo*. With this approach we can recursively reach all references by symbolic references. The expected behavior is that when fetching all references from *refs/<nix-hash>/** that all reachable references will also be fetched. Unfortunately, this is not the case because of a Git inconsistency which has been reported. When fetching symbolic references, Git resolves the symbolic reference to direct references [12]. This makes this approach of fetching package unusable, as the references of the dependencies won't be fetched. But once this upstream bug is fixed, the explained method is a viable approach.

In the second approach there is no *deps* namespace and Gachix fetches each reference explicitly. The initial fetch downloads all Git objects associated with the package closure. The subsequent fetches only download the references. The fetching of the references is done in a breadth-first-search manner.

TODO: should I add the algorithm?

3.3.3. Nix Daemons

If packages don't exist locally and no other replica has the package, it has to be fetched from a Nix store. The Nix daemon protocol can be used to achieve this. With this protocol, we can talk to the local Nix daemon (if one exists) via a socket or to remote Nix daemons via the SSH protocol. Gachix uses the Nix daemon protocol to fetch package metainfo information called "Path Info" in Nix and to download package contents in the NAR format. Gachix parses the NAR files and unpacks them as Git trees or blobs.

4

Implementation

This section explains how a binary cache was implemented using the ideas presented in Section 3. The objective was to create a store for Nix packages using Git and provide them using the common Nix cache interface. The name of this cache is Gachix and the source code is available on Github.³

The project is written in the Rust programming language. This compiled language is ideal for resource-intensive tasks such as parsing a large number of NAR files. The language is also optimal for concurrent tasks, which is used in Gachix for serving multiple connections at once. Rust guarantees memory and thread safety. It eliminates many classes of bugs (e.g. use after free) at compile time.

4.1. Architecture

A high-level overview of the implementation can be seen in Figure 4. The top-level boxes represent the most relevant modules in the code base and the nested boxes their most important functions.

The Command Line Interface (CLI) module gives a friendly interface for interacting with the cache. It also manages the state of the configuration for the binary cache. Gachix can be configured via environment variables or a YAML configuration file. The CLI module merges the configuration options coming from these different sources. With the CLI, the user can start the web server or add a package (and all its dependencies) to the cache.

³<https://github.com/EphraimSiegfried/gachix>

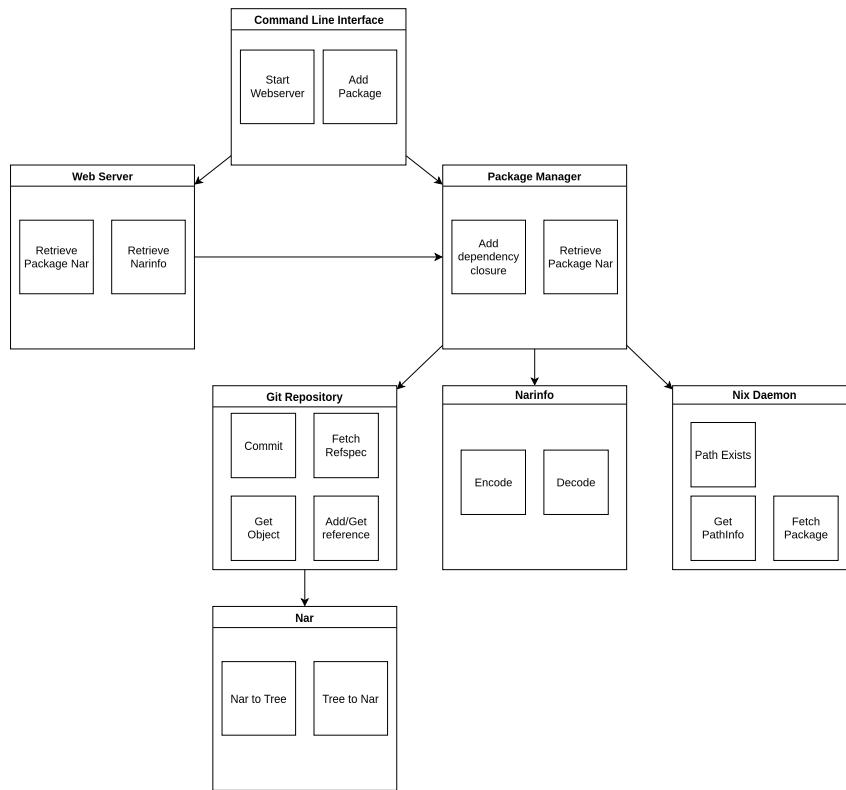


Figure 4: Gachix Architecture Overview

The web server implements the Nix binary interface (See Section 2.1.3.) and serves clients which connect to it via HTTP. It makes requests to the internal package manager module and forwards responses from it to clients.

The package manager module is the core module. It is responsible for doing the Nix to Git mapping discussed in Section 3.1.. It also includes the *add_closure* algorithm discussed in Section 3.3.1..

The Repository module is mainly a wrapper of the Rust library Git²⁴, which is a Rust wrapper of the C++ libgit implementation, which provides low-level Git operations.⁵ The Repository module gives abstractions over common Git operations used for the binary cache, e.g. creating commit objects with constant author, message and date values.

The Nix Daemon module is responsible for communicating to either the local Nix daemon (i.e. the daemon which runs on the same machine as Gachix) or to remote Nix Daemons via the SSH protocol. It contains code for setting up SSH connections, retrieving metadata about store paths and retrieving store objects in the NAR archive format. It depends on a

⁴<https://github.com/rust-lang/git2-rs>

⁵<https://libgit2.org/>

custom fork of a library which implements the Nix daemon protocol.⁶ The fork includes low-level code for retrieving the NAR which did not exist in the original library.

The Narinfo module constructs a Narinfo data structure from the Nix object metadata retrieved from the Nix daemon. It also is able to encode this metadata as a string, which is then stored as a blob in the Git database. Additionally, it signs the signature of the Narinfo and appends it to the narinfo (See Section 2.1.3.).

The NAR module transforms trees to NARs and vice versa. It is used to transform NARs retrieved from Nix daemons to equivalent Git trees. It encodes trees as NARs when Nix cache clients request packages. It does not have to load the whole Git tree onto memory because it is able to stream the NAR, i.e. decode the tree in chunks and serve these chunks continuously.

4.2. Concurrency

To increase the performance of the binary cache, it is crucial to handle requests concurrently. Concurrency can lead to inconsistent state or crashes if handled incorrectly. Inconsistent state happens most often when multiple threads modify the same objects. In Gachix this threat is eliminated by ensuring that threads never modify objects.

Let us consider the two most prevalent operations: Retrieving and adding packages. When a package is retrieved, Gachix looks up the corresponding reference, transforms the referred package tree into a NAR and streams it to the user. The only operations involved in this process are read operations. It does not cause conflict when multiple threads read the same object at the same time.

Adding a package can happen by contacting other Gachix peers and fetching the needed Git objects from them. In this case, only new objects will be added to the Git database. Two threads cannot store two different objects with the same name because trees, blobs and commits are content addressed and references are named after unique Nix hashes. The only scenario that can happen is that two threads try to add the same object but this does not cause a conflict. When fetching NARs from Nix daemons they are transformed to equivalent Git trees. In this process, also only new objects are added.

Since all operations either read or add objects to the database, there is no conflict between multiple threads. Therefore no locking mechanisms have to be used when serving the packages concurrently.

⁶<https://codeberg.org/siegii/gorgon/src/branch/main/nix-daemon>

4.3. Nix Archive

When serving clients, Gachix encodes packages which are stored as Git trees into Nix Archives (Nars). The specification of the Nix Archive Format is shown in Listing 1. [6] The *str* function writes the size of the bytes to be written, the byte sequence specified and a padding of 0s to a multiple of 8 bytes.

```

nar = str("nix-archive-1"), nar-obj;

nar-obj = str("("), nar-obj-inner, str(")");

nar-obj-inner
= str("type"), str("regular") regular
| str("type"), str("symlink") symlink
| str("type"), str("directory") directory
;

regular = [ str("executable"), str("") ], str("contents"),
str(contents);

symlink = str("target"), str(target);

(* side condition: directory entries must be ordered by
their names *)
directory = str("type"), str("directory") { directory-
entry };

directory-entry = str("entry"), str("("), str("name"),
str(name), str("node"), nar-obj, str(")");

```

Listing 1: Backus Naur Form of the Nix Archive

4.3.1. Simple Encoder

A simple approach to encode a given tree into a NAR is to explore the tree in a depth first search manner, retrieve the metadata about the objects and write this in a byte buffer according to the NAR format. Trees are mapped to directories, blobs to symlinks or regulars depending on their type. This approach is fine for shallow trees but is inefficient and can

even lead to errors for large trees. This is because the buffer size can exceed the RAM limit for large packages.

4.3.2. Stream Encoder

A more efficient approach is to stream data chunks to the client, allowing the NAR to be constructed incrementally. To enable this, Gachix implements the Stream trait from the futures crate (A Rust interface) for the encoder module.⁷ Having implemented this trait allows the HTTP server library called Actix Web to asynchronously stream the NAR to the client.⁸ Implementing this trait requires implementing the function `poll_next` which returns the next chunk of bytes to be sent. Recursing the Git tree as in the first approach is not possible anymore because control has to be given back to the caller of `poll_next`.

To address the inability to use standard recursion within the asynchronous `poll_next` method, the implementation replaces the call stack with an explicit state stack (`Vec<TraversalState>`). This transforms the traversal into an iterative state machine. The `TraversalState` enum captures the specific phase of visiting a node, i.e. whether the encoder is initializing a node (`StartNode`), iterating through a directory's children (`ProcessTreeEntries`), or closing a syntactic scope (`FinishNode`). When the encoder encounters a directory, instead of making a blocking recursive call, it simply pushes the corresponding states onto the stack. This effectively schedules the processing of child nodes and the finalize markers for future execution cycles, allowing the function to yield control back to the async runtime after every chunk.

This architecture decouples the logic of tree traversal from the actual transmission of data. As the state machine advances, it serializes specific components of the NAR format such as padded headers, names, or file contents and queues them into a `pending_chunks` buffer. The `poll_next` function prioritizes draining this buffer to the consumer before processing further states from the stack. This ensures that even massive Git trees can be encoded and streamed incrementally without exceeding the server's RAM limits.

4.4. Nix Daemon

In Gachix the Nix daemon protocol is used to interact with the Nix API, i.e. to fetch packages and retrieve metadata information about them. One benefit of using the Nix daemon rather than utilizing the Nix command line interface is that it allows interacting

⁷<https://docs.rs/futures/latest/futures/prelude/trait.Stream.html>

⁸https://docs.rs/actix-web/latest/actix_web/struct.HttpResponseBuilder.html#method.streaming

with the Nix daemon as a non-root user. Another benefit of using it is that it prevents calling system commands such as `nix store dump-path` and `nix path-info`. Spawning a new process for each request to Nix is considerably slower than utilizing inter-process communication to interact with the API through the persistent daemon. Lastly, using the Nix daemon protocol makes it straightforward to interact with remote Nix daemons. Gachix connects to remote nodes via SSH and then calls `nix daemon --stdio`. It then communicates with the remote daemon by piping data through `stdin` and `stdout`.

4.4.1. Libraries

There are few publicly available libraries that implement the Nix daemon protocol interface, and finding them is difficult. The existing implementations identified are typically found as submodules or components within larger projects, rather than as standalone, dedicated libraries. Consequently, these modules are often highly tailored to the specific requirements of their parent projects, often implementing only a limited subset of the full Nix daemon operations necessary for their particular use case. Libraries that were found are:

- **Snix**: Snix reimplements the Nix package manager (more in Section 6.1.). It has a module for interacting with existing Nix protocols. In this module there is a submodule which implements the daemon handshake protocol and a limited set of operations.⁹
- **Harmonia**: Harmonia implements the Nix binary cache interface. It also implements the Nix daemon protocol for the client.¹⁰
- **Gorgon**: Gorgon is a continuous integration framework. It contains a submodule which implements both the client and server daemon protocol.¹¹

Gachix used the Gorgon module because it is the most generic implementation and it has most operations implemented which are needed in Gachix. It only lacks the operation to fetch NARs from Nix stores. This missing feature was implemented in a custom fork.

4.5. Content and Input Addressing Schemes

Nix uses two methods to define the identity of packages. The prevalent way is to use the hash of the input dependency graph, which is stored in the derivation of a package. This type of identity is called input-addressed. The other method is to hash the NAR of a package, which is called content-addressing.

⁹https://git.snix.dev/snix/snix/src/branch/canon/snix/nix-compat/src/nix_daemon

¹⁰<https://github.com/nix-community/harmonia/tree/main/harmonia-daemon>

¹¹<https://codeberg.org/gorgon/gorgon/src/branch/main/nix-daemon>

Gachix uses Git to store packages, which also identifies packages by their content. The difference between Nix and Git is that they use different hashing algorithms and hash different formats. While Nix uses SHA256 and hashes NARs, [11] Git uses SHA1 and hashes its internal Git object format. [5]

There was some effort made to adopt the same hashing scheme as Nix within Gachix. However, this was not finalized because it proved simpler to maintain a mapping between Nix hashes and Git hashes directly within Gachix. This approach offers the benefit of allowing input-addressed Nix hashes to be used for identifying packages in Gachix. Furthermore, the performance and storage overhead of using references, rather than the Git object hashes directly, is negligible.

TODO: I probably have to prove this but I have no idea how

4.6. Limitations

In the current implementation on Gachix a few features are missing that other Nix binary cache services provide. The most notable ones are:

- Normally packages are sent as compressed NARs (most often compressed with `xz`). In the current implementation Gachix only sends plain NARs.
- There is no command implemented which removes packages. However, this is rather easy to implement. It involves removing the reference of the package and calling the Git garbage collector.
- The `.ls` Nix binary cache endpoint is not implemented. This endpoint normally lists the entries of a package.
- Gachix is not capable of signing the Narinfo of a package.

Furthermore, Algorithm 1 (Section 3.3.1.) can be enhanced by implementing a peer-selection heuristic designed to minimize the latency of package addition. The current algorithm is inefficient because it contacts peers in an arbitrary fixed order.

5

Evaluation

This section shows independently verifiable features of Gachix. The following claims about Gachix are made:

1. Gachix is faster in package retrieval than existing cache implementations in 47.4 % of the cases but slower on average. (Section 5.2.)
2. Gachix is more storage efficient than other cache services. (Section 5.3.)
3. Gachix can be deployed on any Unix machine, including on systems without Nix installed. (Section 5.4.)
4. Gachix is transparent to Nix users as it can be used to fetch Nix packages using the Nix substituers interface. (Section 5.5.)

5.1. Functional Comparison to other Cache implementations

There are a few projects which implement the Nix binary cache interface. The most notable ones are:

- **nix-serve**: This is the first cache implementation developed by Eelco Dolstra (i.e. the founder of Nix).¹² It is written in Perl.
- **nix-serve-ng**: This is the successor of nix-serve. It is written in Haskell.¹³
- **harmonia**: This is a modern implementation of the binary cache interface with many features.¹⁴ It is written in Rust.

A notable difference between Gachix and the caches presented above is the other caches directly use the Nix store for storing packages. If a Nix user wants to serve her packages

¹²<https://github.com/edolstra/nix-serve>

¹³<https://github.com/aristanetworks/nix-serve-ng>

¹⁴<https://github.com/nix-community/harmonia>

with Gachix, she has to copy the packages from the Nix store to Gachix. With the other implementations, this is not necessary.

On the other hand, the benefit of using Gachix is that it does not rely on any Nix infrastructure (such as the Nix store) and it can be deployed on a Unix machine without Nix installed. All other implementations expect that Nix is installed on the host machine.

5.2. Package Retrieval Latency

To test whether the retrieval speed of packages is acceptable, Gachix was compared against the cache services presented in Section 5.1..

5.2.1. Methodology

In this benchmark 1000 random packages from the official Nix registry were added to the Nix store and to the Gachix cache. Each cache service was then started and for each package the Narinfo and the NAR was fetched. The end-to-end latency (request sent to full response received) was measured for each request.

5.2.2. Result

The average fetch latency for Narinfo is presented in Figure 5. The average Narinfo retrieval speed is around 0.001 for almost all services except *nix-serve*, which has an average latency of 0.0082 seconds.

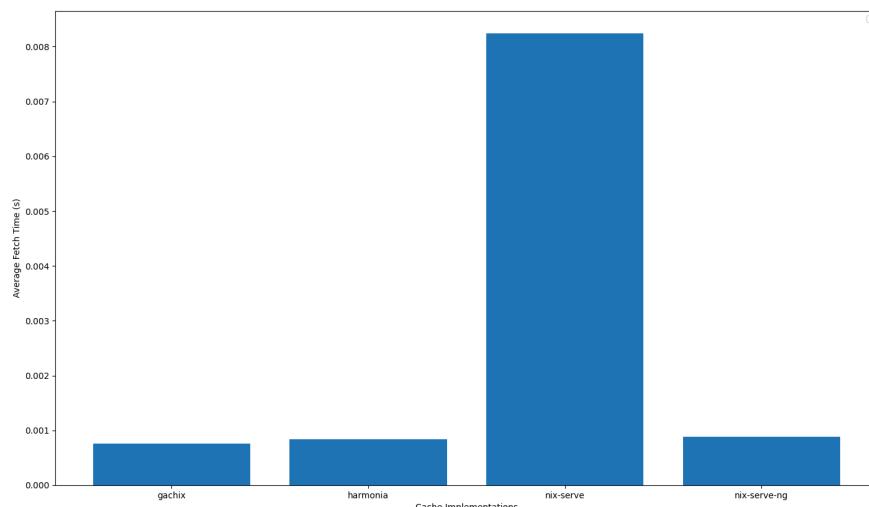


Figure 5: Average Narinfo Fetch Time by Cache Service

The average package fetch time by cache service is shown in Figure 6. With an average latency of 0.01159 seconds *harmonia* is the fastest cache service on average closely followed by *gachix*, which has an average latency of 0.014278 seconds.

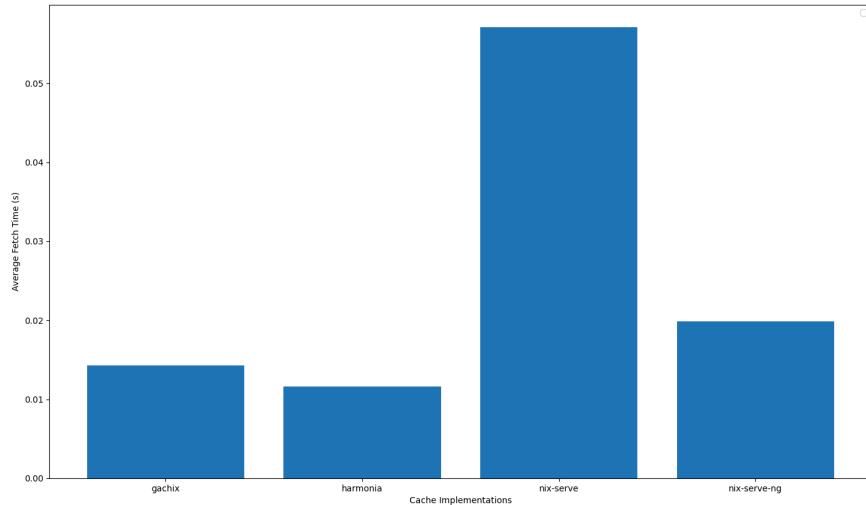


Figure 6: Average Package Fetch Time by Cache Service

The pie chart in Figure 7 shows which services were the fastest among all services. The services *gachix*, *harmonia* and *nix-serve-ng* have an almost equal number of times where they served packages the fastest. Nevertheless *gachix* has shown to be the fastest most times by having been fastest in 474 cases out of 1000.

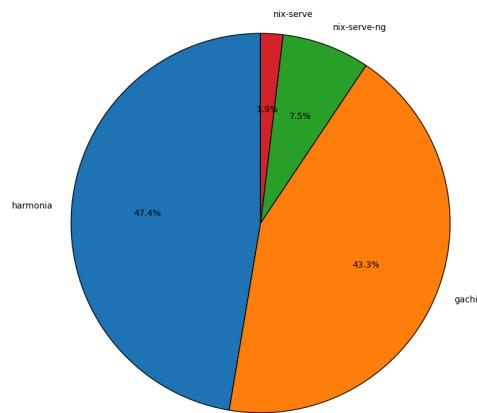


Figure 7: Average Package Fetch Time by Cache Service

5.2.3. Discussion

The reason why *nix-serve* has a much slower latency than the other services is probably because Perl (the language that *nix-serve* was written in) is an interpreted language and all other languages are compiled.

It is interesting that *gachix* performs well in the package retrieval benchmark because it needs to decompress Git objects when constructing the NARs which the other services don't have to because everything in the Nix store is stored as decompressed.

Gachix demonstrates strong performance, achieving a package latency very near the best average and proving to be the fastest in the majority of test cases. It is interesting that *gachix* performs well because it needs to decompress Git objects when constructing the NARs which the other services don't have to because everything in the Nix store is stored as decompressed.

From the results we can conclude that *gachix* is reasonably fast and can compete with other products in this area.

5.3. Package Storage

This benchmark compares the disk storage usage of Gachix to the cache services presented in Section 5.1.. Given that all services use the Nix store and Gachix uses Git as its primary storage for Nix packages, the comparison is more accurately one between the Nix store and the Git database.

5.3.1. Methodology

In this experiment, 1000 randomly selected packages were added to both the Nix store and Gachix.

To assess storage consumption, the total storage used by Gachix was measured by the size of its `.git` directory. This was compared against the sum of the size of all 1000 packages in the Nix store.

Note on Comparison: The sum of the package sizes in the Nix store serves as a lower-bound estimate for the storage required by other cache services. This estimate is conservative because it does not account for potential operational overhead or internal metadata that other caching mechanisms might introduce.

5.3.2. Result

The sum of the package sizes in the Nix store is 3.88 GB. The size of the .git repository is 0.68 GB. This is a size reduction of 82.47%.

5.3.3. Discussion

There are two reasons why we observe this size reduction. Firstly, Gachix compresses its objects using zlib. [11]. The Nix store does not contain any compressed packages.

Secondly, since the Git object database is a Merkle tree and every object is identified by its hash, identical files in the Nix store are only stored once in the Git database. This deduplication of files is also a reason why the size is smaller in Gachix.

5.4. Deployment on any Unix Machine

This section shows that Gachix can be deployed on Unix machines without Nix installed.

5.4.1. Methodology

For this experiment a Dockerfile was written which can be found on the Gachix repository.¹⁵ In this Dockerfile Gachix is built in an environment with Rust installed. After the Gachix binary is built, it is placed in a separate Ubuntu container, where the binary will be run. The experiment proceeded as follows:

- Clone the Gachix source repository
- Build the docker container with `docker build <gachix-repository-path> . -t gachix`
- Run the container with `docker run gachix`

TODO: show that packages can be added to the binary cache, currently it starts the server with 0 packages.

5.4.2. Result

After executing these steps, we can see that Ubuntu container was able to replicate a package from a Nix host. It is also able to serve the package via the Nix binary interface.

¹⁵<https://github.com/EphraimSiegfried/gachix/blob/master/Dockerfile>

5.4.3. Discussion

TODO: what should I write here?

5.5. Nix Transparency

TODO: Needs implementation

5.5.1. Methodology

5.5.2. Result

5.5.3. Discussion

6

Related Work

6.1. Snix

6.2. Laut

7

Conclusion

Bibliography

- [1] E. Dolstra, “The Purely Functional Software Deployment Model,” 2006.
- [2] The NixOs Foundation, “How Nix Works.” Accessed: Sep. 19, 2025. [Online]. Available: <https://nixos.org/guides/how-nix-works>
- [3] Nix Dev, “Store Object.” Accessed: Dec. 16, 2025. [Online]. Available: <https://nix.dev/manual/nix/2.32/store/store-object.html>
- [4] Nix Dev, “Input-addressing derivation outputs.” Accessed: Dec. 16, 2025. [Online]. Available: <https://nix.dev/manual/nix/2.32/store/derivation/outputs/input-address.html#input-addressing-derivation-outputs>
- [5] Nix Dev, “Content-Addressing Store Objects.” Accessed: Dec. 15, 2025. [Online]. Available: <https://nix.dev/manual/nix/2.29/store/store-object/content-address>
- [6] Nix Dev, “Nix Archive (NAR) format.” Accessed: Dec. 04, 2025. [Online]. Available: <https://nix.dev/manual/nix/2.22/protocols/nix-archive>
- [7] F. Zakaria, “Nix Binary Cache.” Accessed: Dec. 04, 2025. [Online]. Available: <https://fzakaria.github.io/nix-http-binary-cache-api-spec/#/>
- [8] Nix Dev <https://github.com/NixOS/nix/blob/5f69fd3e8d99889c571f558113786bd05f0e38ee/src/libstore/ssh-store.cc#L201-L220>, “Multi-User Mode.” Accessed: Dec. 16, 2025. [Online]. Available: <https://nix.dev/manual/nix/2.28/installation/multi-user>
- [9] NixOS, “nix.” Accessed: Dec. 16, 2025. [Online]. Available: <https://github.com/NixOS/nix/blob/5f69fd3e8d99889c571f558113786bd05f0e38ee/src/libstore/ssh-store.cc#L201-L220>
- [10] Snix, “Operations.” Accessed: Dec. 16, 2025. [Online]. Available: <https://snix.dev/docs/reference/nix-daemon-protocol/operations/>
- [11] S. Chacon and Straub Ben, *Pro Git*. Apress, 2025. [Online]. Available: <https://git-scm.com/book/en/v2>
- [12] GitLab and GitLab Organization, “git remote -t' should list tags for specific remotes.” Accessed: Dec. 02, 2025. [Online]. Available: <https://gitlab.com/gitlab-org/git/-/issues/175>